


Two evaluation traps, each invisible to the other's diagnostic

Source-item leakage

rows constructed from one source item, then split by row




exact antisymmetry transfers the memorised score with certainty

$$P(\text{sibling in training}) = 1 - (1 - f)^{m-1} = 0.80 \text{ at } f = 0.8, m = 2$$

Presentation-order prior

the split is clean; the label crosses through the protocol



flip correlation -0.925 — insensitive to a constant offset,
so it reports nothing wrong

fixed-order 0.948 → **strict antisymmetrized 0.639**

a 50% training-time swap protocol did not prevent this;
selecting the checkpoint on fixed-order accuracy banked it